

# **CS 553: Programming Assignment 1**

## **BENCHMARKING Design Document**

**Viral Bhojani**  
**Swapnil Dharawat**

**Introduction:**

Document consist of performance evaluation of 4 different benchmarks which are

1. CPU Benchmarking
2. DISK Benchmarking
3. MEMORY Benchmarking
4. NETWORK Benchmarking

**Consideration:**

Plots are generated using gnuplot. All the experiments were run on OpenStack KVM instances on Chameleon Testbed.

Specifications for the KVM we ran the benchmarks on:

Specification of KVM	
Flavor	m1.medium
Flavor ID	3
RAM	4GB
VCPUs	2 VCPU
Disk	40GB

# Performance Evaluation

## 1. CPU Benchmarking


### 1.1 Decision:

- It is written in C programming language to generate GLOPS and GFLOPS by iterating through 40 operations for each thread (i.e. 1, 2, 4, 8) passed in array and choice for integer and float is also passed in array.
- Strong scaling is achieved as follows. For each number of threads all the 40 operations are executed. So, when 1 thread is passed 40 operations are executed by in thread and when 2 threads are passed, 20 operations are passed in thread 1 and 20 for thread 2.

### 1.2 Report:

Type	Thread 1	Thread 2	Thread 4	Thread 8
<b>GLOPS</b>	1.376806	2.284295	1.843999	1.394237
<b>GFLOPS</b>	1.50988	2.543246	2.03924	1.479079

Graphical Representation:



Theoretical Performance:

$$\begin{aligned} &\text{CPU speed(GHz)} * (\text{Number of CPU cores}) * \text{IPC} \\ &= 3.049 \text{ GHz} * 2 * 2 \\ &= 12.196 \end{aligned}$$

Compared to theoretical performance, achieved is around 20% FLOPS

In Ideal scenario, CPU utilization should be 100%. In our case it was nearly 20% of CPU utilization and hence the difference

### LINPACK benchmarking performance:

```
Number of tests: 15
Number of equations to solve (problem size) : 1000 2000 5000 10000 15000 18000 20000 22000 25000 26000 27000 30000 35000 40000 45000
Leading dimension of array : 1000 2000 5008 10000 15000 18008 20016 22008 25000 26000 27000 30000 35000 40000 45000
Number of trials to run : 4 2 2 2 2 2 2 2 2 2 1 1 1 1 1
Data alignment value (in Kbytes) : 4 4 4 4 4 4 4 4 4 4 1 1 1 1 1
```

Maximum memory requested that can be used=3202964416, at the size=20000

===== Timing linear equation system solver =====

Size	LDA	Align.	Time(s)	GFlops	Residual	Residual(norm)	Check
1000	1000	4	0.070	9.5866	9.394430e-13	3.203742e-02	pass
1000	1000	4	0.016	42.5550	9.394430e-13	3.203742e-02	pass
1000	1000	4	0.012	54.2534	9.394430e-13	3.203742e-02	pass
1000	1000	4	0.012	53.6139	9.394430e-13	3.203742e-02	pass
2000	2000	4	0.082	64.8811	4.085732e-12	3.554086e-02	pass
2000	2000	4	0.081	65.6884	4.085732e-12	3.554086e-02	pass
5000	5008	4	1.160	71.9030	2.262585e-11	3.154992e-02	pass
5000	5008	4	1.163	71.7065	2.262585e-11	3.154992e-02	pass
10000	10000	4	8.948	74.5291	9.187981e-11	3.239775e-02	pass
10000	10000	4	9.258	72.0326	9.187981e-11	3.239775e-02	pass
15000	15000	4	29.693	75.7912	2.219450e-10	3.495671e-02	pass
15000	15000	4	29.499	76.2887	2.219450e-10	3.495671e-02	pass
18000	18008	4	50.912	76.3795	2.886628e-10	3.161212e-02	pass
18000	18008	4	51.704	75.2099	2.886628e-10	3.161212e-02	pass
20000	20016	4	72.156	73.9249	3.669736e-10	3.248520e-02	pass
20000	20016	4	70.404	75.7643	3.669736e-10	3.248520e-02	pass

Performance Summary (GFlops)

Size	LDA	Align.	Average	Maximal
1000	1000	4	40.0022	54.2534
2000	2000	4	65.2848	65.6884
5000	5008	4	71.8048	71.9030
10000	10000	4	73.2808	74.5291
15000	15000	4	76.0400	76.2887
18000	18008	4	75.7947	76.3795
20000	20016	4	74.8446	75.7643

```
[cc@pa1-swapnil-dharawat linpack]$ ./runme_xeon64
This is a SAMPLE run script for SMP LINPACK. Change it to reflect
the correct number of CPUs/threads, problem input files, etc..
./runme_xeon64: line 33: [: too many arguments
Mon Oct 9 23:38:06 UTC 2017
Intel(R) Optimized LINPACK Benchmark data

Current date/time: Mon Oct 9 23:38:07 2017

CPU frequency: 3.049 GHz
Number of CPUs: 2
Number of cores: 2
Number of threads: 2
```

## 2. Memory Benchmarking:

### 2.1 Decisions:

- It is Written in C programming language using multithreading concept for different block sizes (8B, 8KB, 8MB, 80MB) on constant memory size i.e. 1.28 GB and Throughput and latency is calculated in three operations i.e. sequential write(memset), random write(memset) and read+write(memcpy).
- Strong Scaling is achieved by dividing memory size by block size and number of threads.

### Report:


Throughput for Sequential Write, Sequential Read write and Random Write operations for (8B, 8KB, 8MB, 80MB) blocks and Threads 1, 2, 4, 8

### Sequential Write

Block Size	Thread 1	Thread 2	Thread 4	Thread 8
8B	1365.07732	2509.96242	2773.42371	2728.42096
8KB	3726.86497	6602.55517	7285.05588	6914.57604
8MB	3805.47605	7350.8756	7348.60581	7231.31701
80MB	3943.64973	29919.1722	7434.98013	7606.31873

### Graphical Representation

Throughput for Sequential Write

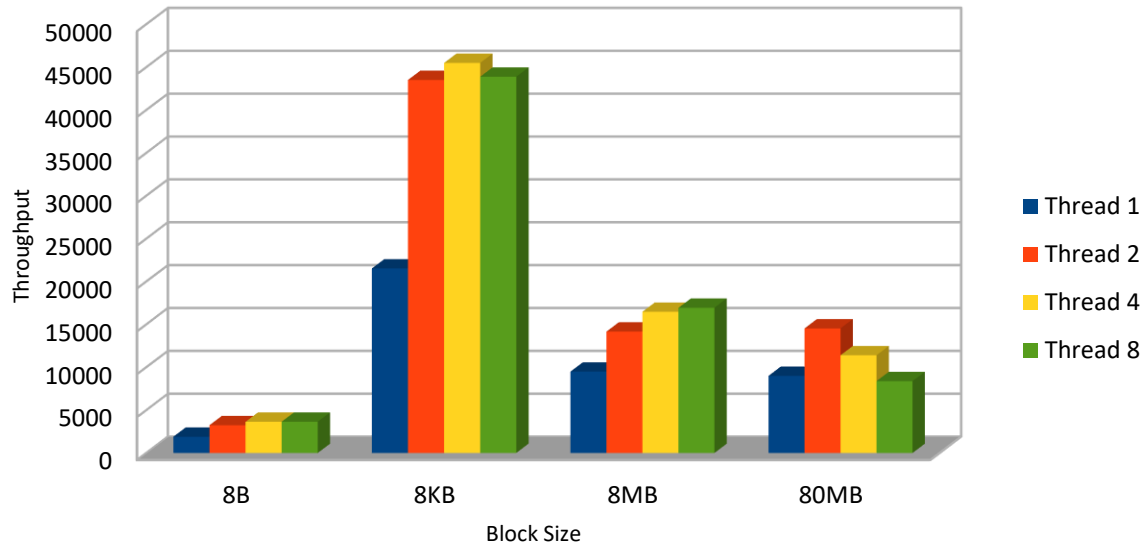


Sequential Read Write

Block Size	Thread 1	Thread 2	Thread 4	Thread 8
8B	1893.37628	3235.33076	3634.64825	3639.76665
8KB	21520.0622	43506.7635	45495.2758	43876.5654
8MB	9486.56303	14153.8984	16463.649	16929.9931
80MB	8992.31248	14524.8224	11393.1001	8376.22311

Graphical Representation:

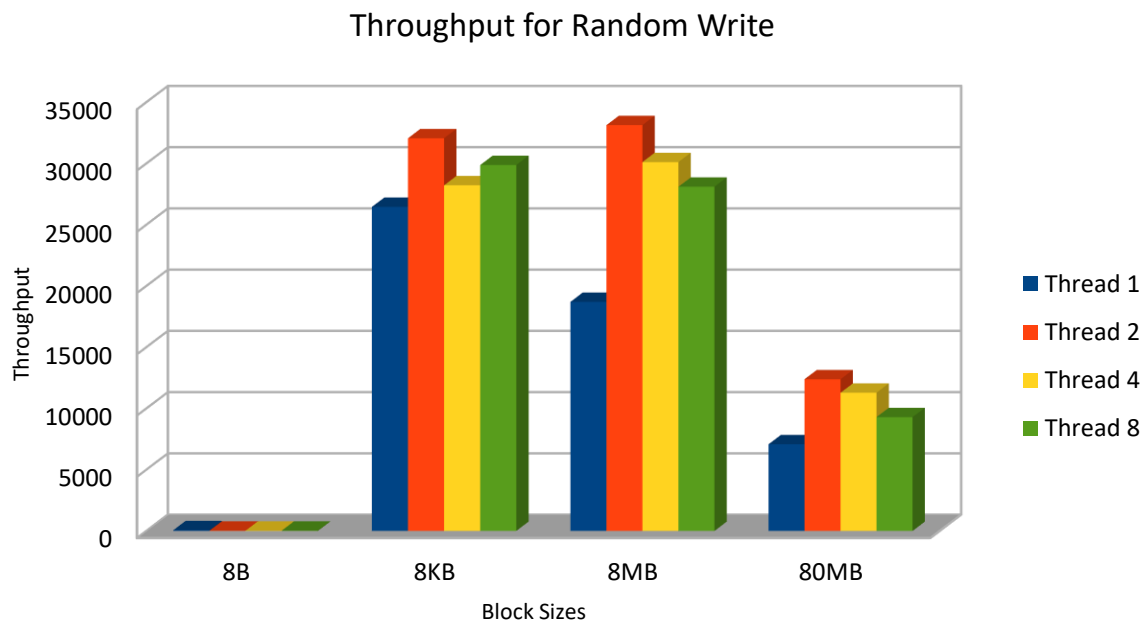
Throughput for Sequential Read Write



Random Write:

Block Size	Thread 1	Thread 2	Thread 4	Thread 8
8B	62.020897	27.470192	25.963327	24.510463
8KB	26459.9538	32061.9019	28220.8267	29876.7344
8MB	18702.6761	33133.3888	30117.0053	28102.8209
80MB	7077.28629	12384.3423	11285.0982	9294.76851

Graphical Representation:




Latency for Sequential Write, Sequential Read write and Random Write operations for (8B) blocks and Threads 1, 2, 4, 8

Sequential Write:

Block Size	Thread 1	Thread 2	Thread 4	Thread 8
8B	5.589E-06	3.0396E-06	2.7509E-06	2.7963E-06

Graphical Representation:






Sequential Read Write:

Block Size	Thread 1	Thread 2	Thread 4	Thread 8
8B	4.0295E-06	2.3582E-06	2.0991E-06	2.0961E-06


Graphical Representation:



Random Write:

Block Size	Thread 1	Thread 2	Thread 4	Thread 8
8B	0.00012301	0.00027773	0.00029385	0.00031127

Graphical Representation:



Theoretical Performance:

**Clock Frequency \* Number of data transfer per clock \* Memory bus interface width \*  
number of interfaces**

## Stream Test

```
[cc@pa1-swapnil-dharawat Memory]$ ls
stream.c
[cc@pa1-swapnil-dharawat Memory]$ gcc stream.c -o stream.out
[cc@pa1-swapnil-dharawat Memory]$ ./stream.out
-----
STREAM version $Revision: 5.10 $
-----
This system uses 8 bytes per array element.
-----
Array size = 10000000 (elements), Offset = 0 (elements)
Memory per array = 76.3 MiB (= 0.1 GiB).
Total memory required = 228.9 MiB (= 0.2 GiB).
Each kernel will be executed 10 times.
The *best* time for each kernel (excluding the first iteration)
will be used to compute the reported bandwidth.
-----
Your clock granularity/precision appears to be 1 microseconds.
Each test below will take on the order of 29272 microseconds.
(= 29272 clock ticks)
Increase the size of the arrays if this shows that
you are not getting at least 20 clock ticks per test.
-----
WARNING -- The above is only a rough guideline.
For best results, please be sure you know the
precision of your system timer.
-----
Function      Best Rate MB/s  Avg time     Min time     Max time
Copy:         6100.6    0.027203     0.026227     0.029107
Scale:        5961.5    0.027666     0.026839     0.029746
Add:          8598.4    0.028981     0.027912     0.031691
Triad:        8064.8    0.030723     0.029759     0.033340
-----
Solution Validates: avg error less than 1.000000e-13 on all three arrays
-----
```

We ran a Stream benchmark and compared all the results obtained. The thing is Stream doesn't give much flexibility to run all experiments with different parameters.

The result when run on Chameleon Testbed is divided into 4 parts having total of 28.05 GB/s.

### 3. Disk Benchmarking:

#### 3.1 Decision:

- It is Written in C programming language using multithreading concept for different block sizes (8B, 8KB, 8MB, 80MB) on constant file size i.e. 10 GB which read and write into the file and Throughput and latency is calculated in three operations i.e. sequential read(memset), random read(memset) and read+write(memcpy).
- Strong Scaling is achieved by dividing memory size by block size and number of threads.


#### 3.2 Report

Throughput for Sequential Read, Sequential Read Write and Random Read operations for (8B, 8KB, 8MB, 80MB) blocks and Threads 1, 2, 4, 8

Sequential Read:

Block Size	Thread 1	Thread 2	Thread 4	Thread 8
8B	16.213795	25.795656	27.089802	28.316206
8KB	5042.983555	7476.984906	8178.652439	7199.829004
8MB	6461.415757	8155.30764	11098.58666	8456.771363
80MB	2681.363893	4743.446262	5686.51319	8317.305193


Graphical Representation:



Sequential Read Write:

Block Size	Thread 1	Thread 2	Thread 4	Thread 8
8B	6.890572	3.067626	4.719271	4.40189
8KB	248.669376	194.182122	770.381248	266.575254
8MB	1253.830059	894.463758	1382.343158	1159.407688
80MB	601.537398	1139.604183	275.645651	2337.067163


Graphical Representation:



Random Read:

Block Size	Thread 1	Thread 2	Thread 4	Thread 8
8B	12.239864	15.55315	17.524179	16.642717
8KB	4289.414863	6254.642117	6442.553063	6691.026184
8MB	5415.170091	8408.329501	9916.254135	8053.404136
80MB	2620.40559	4112.793358	5113.740781	6138.91207

Graphical Representation:




Latency for Sequential Read, Sequential Read Write and Random Read operations for (8B) blocks and Threads 1, 2, 4, 8

Sequential Read:

Block Size	Thread 1	Thread 2	Thread 4	Thread 8
8B	0.000471	0.000296	0.000282	0.000269


Graphical Representation:



Sequential Read Write:

Block Size	Thread 1	Thread 2	Thread 4	Thread 8
8B	0.001107	0.002487	0.001617	0.001733


Graphical Representation:



Random Read:

Block Size	Thread 1	Thread 2	Thread 4	Thread 8
8B	0.000623	0.000491	0.000435	0.000458

Graphical Representation:



ioZone Benchmarking;

The benchmarking is performed on two sets

- i. 8 MB block size and 10GB File size

```
[cc@pai-swapnil-dharawat current]$ ./iozone -a -r 8m -s 10g -T
iozone: Performance Test of File I/O
Version $Revision: 3.471 $
Compiled for 64 bit mode.
Build: linux

Contributors:William Norcott, Don Capps, Isom Crawford, Kirby Collins
Al Slater, Scott Rhine, Mike Wisner, Ken Goss
Steve Landherr, Brad Smith, Mark Kelly, Dr. Alain CYR,
Randy Dunlap, Mark Montague, Dan Million, Gavin Brebner,
Jean-Marc Zucconi, Jeff Blomberg, Benny Halevy, Dave Boone,
Erik Habbinga, Kris Strecker, Walter Wong, Joshua Root,
Fabrice Bacchella, Zhenghua Xue, Qin Li, Darren Sawyer,
Vangel Bojaxhi, Ben England, Vikentsi Lapa,
Alexey Skidanov.

Run began: Mon Oct 9 20:11:03 2017

Auto Mode
Record Size 8192 kB
File size set to 10485760 kB
Command line used: ./iozone -a -r 8m -s 10g -T
Output is in kBytes/sec
Time Resolution = 0.000001 seconds.
Processor cache size set to 1024 kBytes.
Processor cache line size set to 32 bytes.
File stride size set to 17 * record size.

      kB  reflen  write  rewrite  read  reread  random  random  bkwd  record  stride  fwrite  frewrite  fread  freread
10485760  8192  157977  566259  174562  185855  156187  145571  159122  8513316  199345  398161  244121  222982  178622

iozone test complete.
[cc@pai-swapnil-dharawat current]$
```

The performance are as follows:

Sequential read operation is 174562 MBPS, Read+ write 566259 MBPS and Random read is 156187 MBPS

Our performance is 36% of ideal scenario when compared on Random read



ii. 8KB block and 10 GB file Size

```
Contributors:William Norcott, Don Capps, Isom Crawford, Kirby Collins
              Al Slater, Scott Rhine, Mike Wisner, Ken Goss
              Steve Landherr, Brad Smith, Mark Kelly, Dr. Alain CYR,
              Randy Dunlap, Mark Montague, Dan Million, Gavin Brebner,
              Jean-Marc Zucconi, Jeff Blomberg, Benny Halevy, Dave Boone,
              Erik Habbinga, Kris Strecker, Walter Wong, Joshua Root,
              Fabrice Bacchella, Zhenghua Xue, Qln Li, Darren Sawyer,
              Vangel Bojaxhi, Ben England, Vikentsi Lapa,
              Alexey Skidanov.

Run began: Wed Oct  4 06:40:57 2017

Auto Mode
Record Size 8 kB
File size set to 10485760 kB
Command line used: ./iozone -a -r 8k -s 10g -T
Output is in kBytes/sec
Time Resolution = 0.000001 seconds.
Processor cache size set to 1024 kBytes.
Processor cache line size set to 32 bytes.
File stride size set to 17 * record size.

      kB reclen   write  rewrite   read  reread   random  random   bkwd   record   stride   fwrite  frewrite   fread
10485760      8  616447  918950  395315  396752   17583   328404  28619  8390594   24848  827568  616364  428717

iozone test complete.
```

The performance are as follows:

Sequential read operation is 395313 MBPS, Read+ write 918950 MBPS and Random read is 17583 MBPS

Our performance is 34% of ideal scenario when compared on Random read.

#### 4. Network Benchmarking:

##### 4.1 Decisions


It is written in java to perform ping pong operation on both TCP and UDP using multithreading concept on both Client side and Server side to send data on the file upto 64KB from the client using port 5004 and server send response of the same file size.

##### 4.2 Report:

Throughput value for both TCP and UDP for 1, 2, 4, 8 Threads.

Connection Type	Thread 1	Thread 2	Thread 4	Thread 8
TCP	58.99534481	200.4197528	219.6658307	142.2203617
UDP	146.6576426	209.121812	248.8857038	183.6398957


Graphical Representation:



Latency value for both TCP and UDP for 1, 2, 4, 8 Threads.

Connection Type	Thread 1	Thread 2	Thread 4	Thread 8
TCP	1.0589045	0.311698	0.2843885	0.439251
UDP	0.425961	0.2987275	0.25100049	0.340179

## Graphical Representation:



## Iperf Benchmarking:

### Server side:

```
[cc@pa1-swapnil-dharawat Network]$ iperf3 -s
-----
Server listening on 5201
-----
Accepted connection from 192.168.0.121, port 53992
[ 5] local 192.168.0.121 port 5201 connected to 192.168.0.121 port 53992
4
[ ID] Interval           Transfer    Bandwidth
[ 5]  0.00-1.00   sec    3.94 GBytes  33.9 Gbits/sec
[ 5]  1.00-2.00   sec    5.03 GBytes  43.2 Gbits/sec
[ 5]  2.00-3.00   sec    4.93 GBytes  42.4 Gbits/sec
[ 5]  3.00-4.00   sec    4.96 GBytes  42.6 Gbits/sec
[ 5]  4.00-5.00   sec    5.21 GBytes  44.7 Gbits/sec
[ 5]  5.00-6.00   sec    4.94 GBytes  42.4 Gbits/sec
[ 5]  6.00-7.00   sec    4.94 GBytes  42.4 Gbits/sec
[ 5]  7.00-8.00   sec    4.89 GBytes  42.0 Gbits/sec
[ 5]  8.00-9.00   sec    4.98 GBytes  42.8 Gbits/sec
[ 5]  9.00-10.00  sec    4.90 GBytes  42.1 Gbits/sec
[ 5] 10.00-10.04  sec     185 MBytes  43.3 Gbits/sec
-----
[ ID] Interval           Transfer    Bandwidth
[ 5]  0.00-10.04  sec     0.00 Bytes    0.00 bits/sec    send
er
[ 5]  0.00-10.04  sec   48.9 GBytes  41.9 Gbits/sec    re
ceiver
```

## Client side:

```
[cc@pa1-swapnil-dharawat Network]$ iperf3 -c 192.168.0.121 port 22
Connecting to host 192.168.0.121, port 5201
[ 4] local 192.168.0.121 port 53994 connected to 192.168.0.121 port 5201
ID] Interval      Transfer    Bandwidth    Retr  Cwnd
[ 4]  0.00-1.00    sec  4.16 GBytes  35.8 Gbits/sec    0   2.00 MBytes
[ 4]  1.00-2.00    sec  5.02 GBytes  43.1 Gbits/sec    0   2.00 MBytes
[ 4]  2.00-3.00    sec  4.91 GBytes  42.2 Gbits/sec    0   2.00 MBytes
[ 4]  3.00-4.00    sec  4.99 GBytes  42.8 Gbits/sec    0   2.00 MBytes
[ 4]  4.00-5.00    sec  5.20 GBytes  44.7 Gbits/sec    0   2.00 MBytes
[ 4]  5.00-6.00    sec  4.93 GBytes  42.4 Gbits/sec    0   2.00 MBytes
[ 4]  6.00-7.00    sec  4.98 GBytes  42.8 Gbits/sec    0   2.00 MBytes
[ 4]  7.00-8.00    sec  4.90 GBytes  42.1 Gbits/sec    0   2.00 MBytes
[ 4]  8.00-9.00    sec  4.89 GBytes  42.0 Gbits/sec    0   2.00 MBytes
[ 4]  9.00-10.00   sec  4.95 GBytes  42.5 Gbits/sec    0   2.00 MBytes
-- -- -- -- --
ID] Interval      Transfer    Bandwidth    Retr
[ 4]  0.00-10.00   sec  48.9 GBytes  42.0 Gbits/sec    0
sender
```